

Random Forests - To predict Whether Patients has diabetics or not

Random forest is a commonly-used machine learning algorithm trademarked by Leo Breiman and Adele Cutler, which combines the output of multiple decision trees to reach a single result. Its ease of use and flexibility have fueled its adoption, as it handles both classification and regression problems.

In [1]:

```
#Setup
import warnings
warnings.filterwarnings("ignore")
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, precision_recall_curve, roc
from sklearn.ensemble import RandomForestClassifier
```

In [6]:

```
names = ["times_pregnant", "glucose_tolerance_test", "blood_pressure", "s
         "bmi", "pedigree_function", "age", "has_diabetes"]
data = pd.read_csv('data/pima-indians-diabetes.csv', names=names, header=
```

In [7]:

```
print(data.shape)
```

(767, 9)

In [8]:

```
#Print no of integers, floats and strings  
data.dtypes.value_counts()
```

Out[8]:

```
int64      7  
float64    2  
dtype: int64
```

In [9]:

```
#Data should be numerical  
data.head()
```

Out[9]:

	times_pregnant	glucose_tolerance_test	blood_pressure	skin_thickness	insulin
0	1	85	66	29	0
1	8	183	64	0	0
2	1	89	66	23	94
3	0	137	40	35	168
4	5	116	74	0	0

In [10]:

```
#Print no of entries for each color  
data.has_diabetes.value_counts()
```

Out[10]:

```
0      500  
1      267  
Name: has_diabetes, dtype: int64
```

In [11]:

```
#Print % of each colors  
data.has_diabetes.value_counts(normalize=True)
```

Out[11]:

```
0    0.65189  
1    0.34811  
Name: has_diabetes, dtype: float64
```

Preprocessing Steps

1. Select Features and Split to X and y .

In [12]:

```
X = data.iloc[:, :-1].values  
y = data["has_diabetes"].values
```

2. Split data to train and test .

In [13]:

```
# Split the data to Train, and Test (75%, 25%)  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
```

Modelling with Random Forest

1. Train a Random Forest model with 200 trees on the training data.

In [14]:

```
## Train the RandomForest Model
rf_model = RandomForestClassifier(n_estimators=200)
rf_model.fit(X_train, y_train)
```

Out[14]:

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_
weight=None,
                        criterion='gini', max_depth=None, max
_features='auto',
                        max_leaf_nodes=None, max_samples=Non
e,
                        min_impurity_decrease=0.0, min_impuri
ty_split=None,
                        min_samples_leaf=1, min_samples_split
=2,
                        min_weight_fraction_leaf=0.0, n_estim
ators=200,
                        n_jobs=None, oob_score=False, random_
state=None,
                        verbose=0, warm_start=False)
```

2. Calculate the accuracy and roc_auc_score of the predictions.

In [15]:


```
# Make predictions on the test set
y_pred_class_rf = rf_model.predict(X_test)
y_pred_prob_rf = rf_model.predict_proba(X_test)

print('accuracy is {:.3f}'.format(accuracy_score(y_test, y_pred_class_rf)))
print('roc-auc is {:.3f}'.format(roc_auc_score(y_test, y_pred_prob_rf[:,1])
```

```
accuracy is 0.766
roc-auc is 0.799
```

In [16]:

```
def plot_roc(y_test, y_pred, model_name):  
    fpr, tpr, thr = roc_curve(y_test, y_pred)  
    fig, ax = plt.subplots(figsize=(8, 8))  
    ax.plot(fpr, tpr, 'k-')  
    ax.plot([0, 1], [0, 1], 'k--', linewidth=.5) # roc curve for random  
    ax.grid(True)  
    ax.set(title='ROC Curve for {} on PIMA diabetes problem'.format(model_name),  
          xlim=[-0.01, 1.01], ylim=[-0.01, 1.01])  
plot_roc(y_test, y_pred_prob_rf[:, 1], 'RF')
```



Receiver Operating Characteristic(ROC) curve is a plot of the true positive rate against the false positive rate. It shows the tradeoff between sensitivity and specificity.

In []: